IAE PARIS — SORBONNE BUSINESS SCHOOL · IDGP ACADEMY

Lean Six Sigma & Continuous Improvement

Reading notes for the morning session. Module 3 — Methods, Tools and Operational Efficiency, Session 4.

17 October 2026 · 09:00–12:00

Prof. Dr. Thiago Lima

International Certificate in Project Management and Organizational Performance

No preparation is required. These notes carry the definitions, figures and sources, so the time in the room can be spent on judgement rather than on terminology.

Contents

The session builds one argument in eleven movements: value is defined from the customer backwards, waste is what the customer would not pay for, and the efficiency that matters is the unit's journey rather than the resource's occupancy.

ONE	The puzzle	4
TWO	Value, seen from the customer backwards	5
THREE	The anatomy of waste	9
FOUR	The efficiency paradox	12
FIVE	The Toyota Production System, as a system	17
SIX	Lean in the language of capital	20
SEVEN	DMAIC — conducting improvement, analysing problems	27
EIGHT	The third equation	30
NINE	Group activity — The Value Case	32
TEN	Going further — how capability accumulates	38
ELEVEN	Sources	41

> **Bring three things.** One recurring process that takes longer than it should; two rough figures — projects in flight, and projects closed in the last twelve months; and one authorised project you know well enough to defend, which is the raw material for the group activity in Part Nine.

Most organisations that adopt lean copy the tools. Few reproduce the results. These notes build the argument for why — and the argument runs on three equations rather than on a list of techniques.

They open with a puzzle. **The answer is deliberately not in this document.** You will be asked to write your hypothesis down in the first ten minutes of the session, and we will test it against the framework at the end of the morning — so if you are reading ahead, stop at the end of Part One and let the morning do the rest.

> **WHAT THE SESSION WILL ASK OF YOU**
>
> Nothing to prepare, but three things to have in mind — all of them about your own organisation, all answerable from what you already know.
>
> **One recurring process** you own or depend on, that takes longer than it should. Contract approval, project initiation, monthly close, incident resolution.
>
> **Two rough figures:** how many projects are in flight in your unit, and how many closed in the last twelve months. Approximations are fine.
>
> **One authorised project** you know well enough to defend. It closes the session as the group activity in Part Nine. A project you have private doubts about is better material than one you are proud of.

A note on framing. The examples below run through factories more often than through project offices, because that is where the ideas were developed and tested. They were not developed for factories. Each section ends with a **bridge to project management**, because the mechanisms transfer exactly — queues are queues, variability is variability — and because this programme is about delivery rather than manufacturing.

ONE

The puzzle

Fremont, California

1982. General Motors closes its Fremont assembly plant. Absenteeism runs near 20%. The workforce is represented by what GM managers describe as the most combative union local in the system. Quality is the worst in General Motors.

1984. The plant reopens as NUMMI, a joint venture with Toyota. Substantially the same workforce is rehired, including the union leadership. Same building. Same equipment. Within roughly two years, quality and productivity approach Toyota's Japanese levels and lead the entire GM system.

Same assets. Same people. **What changed — and where would that change appear on a balance sheet?**
And a second question, harder: GM had access to NUMMI from the inside for two decades. Why was it never replicated across the rest of the network?

Hold your answer. Everything that follows is a way of making it precise.

TWO

Value, seen from the customer backwards

Taiichi Ohno described the whole of the Toyota Production System in one sentence: *"All we are doing is looking at the timeline — from the moment the customer gives us an order to the point when we collect the cash — and we are reducing that timeline by removing the non-value-added wastes."*

That sentence does two things at once. It defines the unit of analysis as **elapsed time from order to cash**, not cost per unit. And it makes the customer, not the accounting system, the judge of what counts as work.

Three kinds of activity

Type	Test	Typical share
Value-adding	Transforms the product or service in a way the customer would knowingly pay for	Under 5%
Necessary non-value-adding	The customer would not pay for it, but law, technology or contract requires it	10–30%
Waste (*muda*)	Neither	The remainder

In value streams that have never been deliberately improved, **value-adding time is a low single-digit percentage of elapsed time.** Everything else is queue. This is the single most consequential number in operations, and almost no organisation measures it.

What the customer is actually choosing

"What the customer would pay for" is not a marketing abstraction. It decomposes into operational variables.

The **competitive priorities** — cost, quality, delivery speed, flexibility — are the operational determinants of willingness to pay. Terry Hill's distinction then sorts them by function. An **order qualifier** is a threshold: its presence attracts nobody, its absence eliminates you from consideration. An **order winner** is what makes a customer choose you once you are in the set. Nobody selects a restaurant because its bathroom is excellent; plenty never return because it was dirty.

The managerial consequence is uncomfortable. Effort poured into qualifiers above the threshold consumes capacity and produces no preference. To be excellent at an order winner, an organisation has to accept being merely adequate elsewhere — and a firm that is mid-table on every attribute has made no choice at all.

David Garvin's eight dimensions of quality sharpen this further, because quality is not one variable. **Conformance** lowers cost; **reliability, durability and serviceability** raise willingness to pay. The same quality effort therefore pays twice, on two different edges of value.

IN PROJECT TERMS

The same distinction governs portfolio selection. Most scoring models treat every criterion as additive — strategic fit, NPV, risk, resource availability — which quietly assumes they are all order winners. They are not. Compliance and regulatory projects are qualifiers: they buy no advantage, and not doing them ends the business. Growth and differentiation projects are winners.

A portfolio that scores both on the same scale will systematically over-fund qualifiers, because they are easier to justify and harder to argue against. The discipline is to fund qualifiers to the threshold and no further, and to concentrate the remaining capacity where preference is actually created.

The value stick

Brandenburger and Stuart formalised where value comes from:

$$\text{Value created} = \text{WTP} - \text{SOC}$$

willingness to pay, minus supplier opportunity cost

Between those two bounds sit three claims: customer surplus (WTP − price), firm margin (price − cost), and supplier surplus (cost − SOC). Negotiating over price or over supplier terms moves value *between* parties. Only two moves **create** it: push WTP up, or push SOC down.

Lean is unusual because it does both at once. Removing waste lowers cost; the same work raises reliability and speed, which raise what the customer will pay. That is the whole strategic claim of the method, and the rest of these notes explain the mechanism.

From value created to advantage

Creating value and having an advantage are different things, and the value stick alone cannot tell them apart. A firm that raises WTP and lowers cost in a market where every competitor is doing the same has created value and captured none of the difference. Advantage is a *comparison*, so it has to be written as one:

$$VC = (WTP_A - Cost_A) - (WTP_B - Cost_B)$$

the value firm A creates, minus the value created by its relevant rival B

Competitive advantage is the **difference between the value two firms create** — not the value one of them creates in isolation. Three consequences follow, and each of them changes how an improvement programme

should be argued.

- **The benchmark is a rival, not last year.** An improvement that cuts your own cost by eight per cent while the reference competitor cuts theirs by twelve has made the firm better and its position worse. Internal time series cannot detect this; only the comparison can.
- **Both terms can be attacked.** VC rises when A improves and equally when B's advantage is made harder to sustain. Most of the literature in these notes is about the first term. The second is why capability that resists imitation matters more than capability that does not.
- **Operational effectiveness erodes by construction.** If the improvement is visible and copyable, B eventually makes the same move, the two parentheses converge, and VC tends to zero. The firm is left creating more value than before and capturing no more of it than its rival. This is Porter's objection, stated as arithmetic rather than as an opinion, and the afternoon session returns to it as the central risk of benchmarking.

The practical use of the formula is in the third reading. It gives a single test to put to any improvement project: *when the competitor does this too, what is left?* If the answer is nothing, the project is a qualifier and should be funded to the threshold and no further. If something is left, name it — that residue is the advantage.

Where the residue can live

Naming it is easier with the identity that organises the rest of the day. Van Mieghem and Allon describe an operation through four linked elements — value, capabilities, assets and processes — and the relationship between them answers the question VC leaves open.

$$V = C \times (A + P)$$

Lima (2026), Ch. 1 — after Van Mieghem & Allon (2015)

A is what the organisation owns: plants, machines, systems. **P** is how work flows across them. **C** is its repeated ability to improve P. The form of the expression is the claim: assets and processes are *additive and purchasable*, capability is a *coefficient*.

Read against VC, that settles where the difference between two firms can survive. Anything in the bracket is available to a competitor at roughly the same price, so improvements that live there move both parentheses and the difference returns towards zero. Only the coefficient resists, because it is the one term nobody can invoice you for. **This is the whole argument of the day, stated early so that everything after it can be read as evidence for or against it** — and Part Eight returns to the identity once there is enough material to test it.

The five principles, and where each one is treated

Womack and Jones reduced the method to five principles in sequence. They are worth stating as a set once, because the sequence is itself the argument: every principle after the first is meaningless until the first has been answered.

	Principle	What it asks	Where
1	**Value**	What is the customer actually paying for? Defined from the customer backwards, never from the process forwards.	This part
2	**Value stream**	Which activities deliver that value, which merely enable it, and which do neither? The map is the instrument.	Part Three
3	**Flow**	Can the work move without stopping? Stoppage, not slowness, is what makes lead time long.	Part Four
4	**Pull**	Is work released by downstream demand, or pushed by upstream capacity and a plan?	Part Five
5	**Perfection**	Does the cycle repeat, with the standard rewritten each time? This is the principle that decides whether the first four survive.	Parts Seven and Ten

THREE

The anatomy of waste

Ohno catalogued **seven** wastes, and that is the canonical list; modern practice adds an eighth, unused talent, which is not Ohno's and is the one most relevant to knowledge work. The list matters less than its structure: **one of them is the root, and the others are symptoms.**

Overproduction is the root waste

Producing more than is needed, earlier than it is needed, generates the other six. It creates inventory, which needs transport and motion, which hides defects, which forces waiting when the defect is finally found. Attack overproduction and the others shrink; attack the others individually and overproduction regenerates them.

Waste	In a factory	In project and knowledge work
Overproduction	Making more than demand	Reports nobody reads, features nobody uses, gold-plating
Waiting	Idle machines and people	Approval queues, decision latency, blocked dependencies
Transport	Moving materials	Handoffs between departments and systems
Over-processing	Machining beyond specification	Sign-off chains, redundant status meetings, precision beyond need
Inventory	Stockpiles and work in process	Half-finished work packages, open tickets, projects started but unstaffed
Motion	Unnecessary movement of people	Searching for information, tool switching, duplicated data entry
Defects	Scrap and rework	Rework loops, failed deliverables, change requests from unclear requirements
Unused talent	Ignored operator ideas	Experts doing administration; front-line insight never reaching decisions

Making it visible — the value stream map

The list above is a vocabulary. A value stream map is the instrument that locates the items in a specific process, and it is the second Lean principle made operational. It is drawn by walking the process, not by assembling it from what people say in a meeting — which is why the first draft almost always contradicts the official procedure.

A map carries three things at once, and the third is the one that does the work.

Layer	What it records	What it reveals
The work	Each step, in sequence, with who performs it	Handoffs — which are transport, and where the defects enter
The information	What triggers each step: a signal, a schedule, a queue, an approval	Whether the stream is pulled or pushed, and where decisions wait on data
The timeline	For each step, the time actually worked and the time elapsed before the next step starts	The ratio between the two, which is flow efficiency, and is the number the map exists to produce

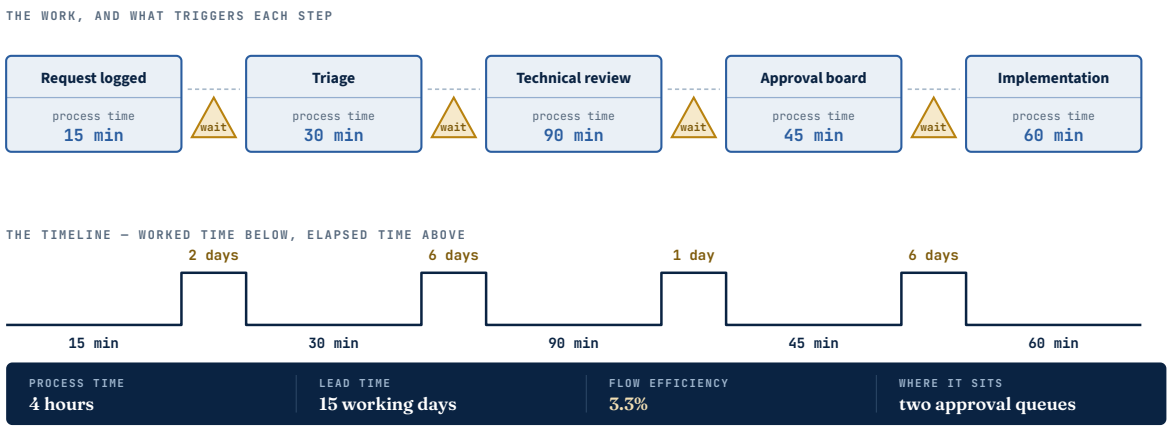


Figure 1. A value stream map of an administrative process — a change request from arrival to implementation. The boxes carry the worked time, the triangles the waiting. The staircase at the bottom is the whole point of the exercise: four hours of work spread across fifteen working days, with almost all of the elapsed time sitting in two approval queues.

Two figures come out of the timeline. **Process time** is the total time anyone is actually working on the unit. **Lead time** is the elapsed time from the first step to the last. In administrative and project processes the first is typically between two and ten per cent of the second, and the gap is not idleness — it is queue. Everything in the waste table lives in that gap.

A useful map adds one more figure per step: the proportion of items arriving **complete and accurate**, needing no clarification or rework. A chain of steps each at 90% complete-and-accurate does not deliver 90%, and the arithmetic of that is in Part Four.

> **CURRENT STATE, THEN FUTURE STATE**
>
> The map of today is diagnosis. The second map — the future state — is the improvement plan, and it is what makes the exercise more than documentation. The discipline is that every difference between the two maps must be attributable to a named change with an owner. A future state nobody can trace back to specific interventions is a wish drawn in boxes.

IN PROJECT TERMS

The process worth mapping in a project organisation is rarely the delivery work itself. It is the surrounding administrative stream: the path from a change request to an approved change, from an invoice to a payment, from a defect report to a closed defect. These are repeatable, they are measurable, and they are where the elapsed time actually accumulates.

Map one of them end to end and the characteristic finding appears: four hours of work spread across three weeks, with the three weeks sitting in two approval queues that exist because of an authority limit nobody has revisited since it was set.

Muda is the symptom; mura and muri are the disease

Three Japanese terms, and executives usually learn only the first.

Mura — unevenness

Irregular flow of work. Unevenness creates queues, and queues create the safety buffers that look like prudent management and are in fact inventory.

Muri — overburden

People and equipment pushed beyond reasonable limits. Overburden creates defects, breakdowns and absence — which then create more unevenness.

This is why improvement programmes plateau after year one. The first wave removes visible muda and produces real gains. Then the gains stop, because the unevenness and overburden that generate muda were never touched, and they regenerate it.

IN PROJECT TERMS

Chronic over-allocation is muri with a resource-management label on it. A specialist assigned at 120% across four projects produces the entire symptom set: defects, because attention is divided; waiting, because every project queues behind the other three; and unused talent, because the scarce expert spends the day reconciling status rather than solving problems.

Mura shows up as the quarter-end delivery peak that every organisation knows about and none levels. The peak is treated as a fact of the calendar; it is usually a consequence of how work was released.

FOUR

The efficiency paradox

Lean is routinely introduced as "doing more with less". Taken literally, that is the definition of the thing lean exists to correct.

Resource efficiency

Keep every person and every machine busy. It optimises what cost accounting measures, which is why organisations default to it.

It generates queues, rework and expediting — and it destroys performance globally while every local unit looks efficient.

Flow efficiency

Optimise the unit's journey through the system. It measures the order, not the resource.

It usually improves both once queues collapse — and almost nobody reports it.

Flow efficiency is simply the share of elapsed time in which the unit is actually being worked on — process time over lead time, the ratio the value stream map in Part Three was built to produce. The figure is usually a shock the first time it is computed, and the shock is the point.

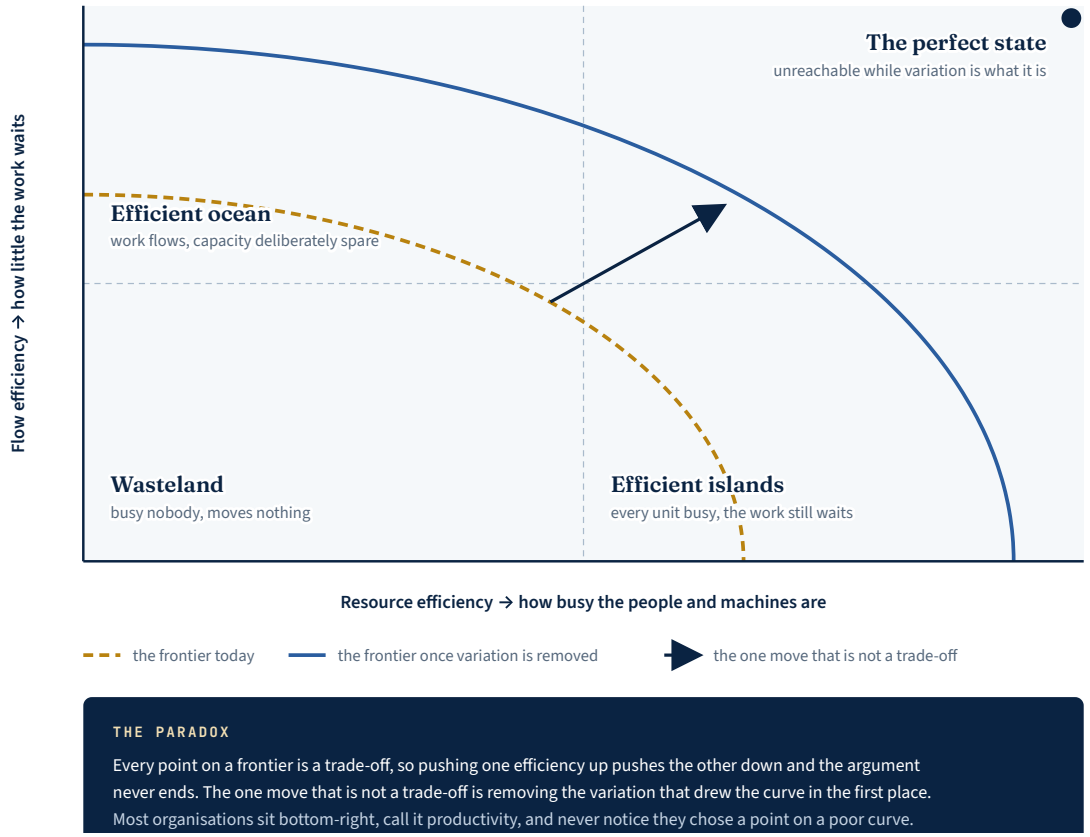


Figure 2. The efficiency matrix, after Modig and Åhlström (2012). Resource efficiency and flow efficiency are not two names for the same thing, and the curve is the reason they trade off: with variation at its current level, the combinations beyond the frontier cannot be bought with effort. Removing variation is what moves the frontier towards the corner where both are high.

The physics of lead time

The relationship between utilisation and queue time is not a management opinion. It is a property of any system with variability, and it is steep:

- As utilisation rises toward 100%, **queue time rises towards infinity**. The curve is flat until roughly 80% and then climbs violently.
- Therefore the instinct to keep every resource busy is **the direct mechanical cause of the inventory it was meant to avoid.**
- Raising utilisation does not raise the system's output. It raises the queue in front of the system.

> **THE SENTENCE WORTH MEMORISING**
>
> **Variability must be absorbed by inventory, by capacity, or by time.** Every organisation chooses one — usually without deciding to. Only one of the three is invisible in the P&L, which is why it is the one most organisations choose by default.

The other half of the method — variation

Everything above is about variability without naming its source, and this is where the Six Sigma half of Lean Six Sigma enters. Lean attacks the time the work spends not being worked on. Six Sigma attacks the **dispersion** of the process — and dispersion is precisely what makes lead time unpredictable. They are not alternative programmes. They are two attacks on the same quantity, which is why the integration is more than a brand.

Two kinds of variation, and only one is a problem to be solved

Shewhart's distinction governs everything that follows. **Common-cause** variation is inherent in the process as designed: it is the ordinary scatter a stable process produces, and no amount of investigating individual cases will reduce it, because every case has a different story and none of the stories is the cause. **Special-cause** variation comes from something that was not there before, and it can be found and removed.

The managerial consequence is uncomfortable and expensive to ignore. **Reacting to common-cause variation as though it were special cause makes the process worse**, because each reaction is an unplanned adjustment to a stable system. The monthly meeting that demands an explanation for every variance against target, favourable or adverse, is this error institutionalised — and it is the single most common misuse of a performance dashboard.



Figure 3. A control chart of the same process over twenty-four measurements. The limits are calculated from the process itself, not from the target. Two signals are present and everything else is noise — and the costly managerial error is not missing the two signals, it is explaining the noise.

Capability — the promise is the spread, not the average

A process has a voice, which is its natural spread, and the customer has a requirement, which is a tolerance. Capability compares the two.

$$C_p = (USL - LSL) / 6\sigma$$

the width the customer allows, over the width the process produces

C_p assumes the process is centred on the target. C_{pk} removes that assumption and is therefore the one that matters, because a process can be admirably narrow and reliably in the wrong place. A conventional threshold is $C_{pk} \geq 1.33$, and the operational reading of it is simple: **the average tells you almost nothing about what the customer will experience; the spread, relative to the tolerance, tells you everything.**

This transfers directly to services and to projects, where the tolerance is usually a promise rather than an engineering specification. A response committed at 48 hours with an average of 30 and a spread that reaches 90 is a process that will break its promise regularly while reporting comfortable averages every month.

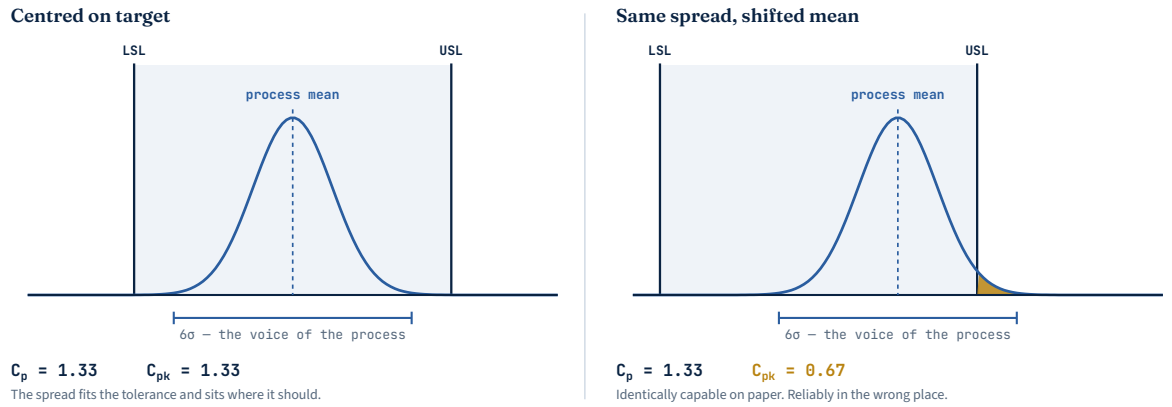


Figure 4. Two processes with identical spread and opposite fates. On the left the mean sits on target and the tolerance is comfortable. On the right nothing about the process changed except where it is centred — C_p is unmoved, C_{pk} halves, and the shaded tail is the share of output the customer will reject. This is why the average is a poor promise and the spread, read against the tolerance, is the real one.

What a sigma level actually says

Sigma level	Defects per million	Right first time	What it means in practice
3σ	66,807	93.3%	Typical of a process nobody has studied. Feels acceptable in isolation.
4σ	6,210	99.38%	A well-run process with active controls.
5σ	233	99.977%	Requires design changes, not vigilance.

6σ	3.4	99.99966%	The benchmark the method is named after.

These figures follow the convention of allowing a 1.5σ long-term drift in the process mean, which is why 6σ corresponds to 3.4 defects per million rather than to the two-per-billion a static normal distribution would give. It is a convention, not a law, and it is worth knowing that it is one.

THE ARITHMETIC THAT MAKES THE CASE TO AN EXECUTIVE

Quality is almost never argued step by step; it is argued at the step and experienced across the chain. A twenty-step process in which **every step is 99% right** delivers a correct outcome **81.8%** of the time. At 99.38% per step — a 4σ process — it delivers 88.3%. At 99.9% per step it delivers 98.0%.

So the familiar objection, *our quality is already good*, is usually true at the step and false at the outcome. It also explains why removing steps is a quality intervention and not only a cost one: every handoff you delete is a term removed from that product.

WHY THE TWO HALVES NEED EACH OTHER

Lean without Six Sigma removes steps and queues and then cannot hold the gain, because the variation that generated the queues is still there and rebuilds them as buffers. Six Sigma without Lean produces beautifully capable steps inside a process whose elapsed time is still ninety per cent waiting — a stable process delivering slowly and predictably.

The connection back to Part Two is direct. Reducing variation lowers cost through less rework, which moves the SOC edge; and it raises reliability and conformance, which moves the WTP edge. It is the same double movement the value stick predicted, now with a measurable instrument attached.

IN PROJECT TERMS

The three buffers have exact project equivalents: **inventory** becomes work in progress and partially completed deliverables; **capacity** becomes overtime, contractors and resource reserve; **time** becomes schedule contingency. Goldratt's critical chain is the argument that these buffers should be sized deliberately and placed where they protect the system, rather than hidden inside every individual task estimate — where they are consumed silently and protect nothing.

The utilisation point lands hard here. A project organisation that plans its people at 100% allocation has built a system in which every piece of variability — one illness, one late approval, one changed requirement — is absorbed by schedule. The resulting delay is not poor project management. It is the planning assumption, working exactly as specified.

FIVE

The Toyota Production System, as a system

The TPS is conventionally drawn as a house, and the drawing is an argument: a roof needs pillars, pillars need a foundation, and removing any element collapses the rest.

Goals — highest quality · lowest cost · shortest lead time · best safety · high morale

Just-In-Time	People and teamwork	Jidoka
Takt time · continuous flow · pull and WIP limits · SMED	Kaizen · *genchi genbutsu* · structured problem solving. The centre of the house, and the multiplier.	Autonomation · stop-and-fix · andon · poka-yoke

Stability — *heijunka* (levelling) · standardised work · 4M stability

The sequence is physics, not preference. Pull cannot function on an unstable, unlevelled base: it either starves or floods. This is why organisations that install kanban boards onto chaotic processes get chaos with better visualisation.

The left pillar: Just-In-Time

Three mechanisms, and one widespread misconception to clear first.

> PULL IS NOT JIT, AND NEITHER IS KANBAN
>
> **A pull system is a system that explicitly limits work in process.** That is the definition. A push system releases work according to an external schedule. Kanban is merely one mechanism for limiting WIP; CONWIP is another. Treating the cards as the essence of the system is the most common conceptual error in the boardroom (Hopp & Spearman, 2004).

- **Continuous flow** — remove the interruptions, batching and breakdowns that convert value-adding time into elapsed time.
- **Pull** — produce nothing until a downstream signal indicates demand. The signal is what stops overproduction.
- **SMED** — batch sizes exist because changeovers are expensive. Kill the changeover and the economic batch size collapses towards one.

Underneath all three sits *heijunka*: levelling the schedule so that the pull signal is stable enough to be worth following.

The right pillar: Jidoka

Quality built in at the source, through three mechanisms of increasing cost and decreasing frequency:

- **Poka-yoke** — design the process so the mistake physically cannot be made.
- **Andon** — signal the abnormality the instant it occurs, not at the end of the shift.
- **Stop-and-fix** — halt the line so that the defect never travels downstream.

Stopping the line looks like the most expensive thing a factory can do, and that reading misses the trade. **A problem that stops the line gets solved; a problem that is worked around does not.** The cost of stopping is visible and immediate; the cost of working around is invisible and permanent.

> IN PROJECT TERMS
>
> Three elements of the house are already in your vocabulary under other names. *Heijunka* is resource levelling. *Standardised work* is the methodology, templates and checklists the PMO maintains. *Poka-yoke* is the mandatory field that will not let the form be submitted incomplete.
>
> The one with no comfortable equivalent is the andon cord. In most project organisations, status reporting is the opposite of an andon: it aggregates upward on a monthly cycle, and a red status costs the reporter more than it costs the organisation. That is precisely a system in which abnormalities are signalled late and quietly. The question to carry into the session is not whether your governance has escalation paths written down — it is how long someone knew, the last time something went seriously wrong.

Why the system beats the tools

This is the best-evidenced claim in the field, and it is the reason the house is drawn as a house rather than as a menu.

Source	Finding
Milgrom & Roberts (1990)	Modern manufacturing is a system of *complementarities*: the return to any one practice rises with the adoption of the others. Partial adoption yields near-zero or negative returns.
MacDuffie (1995)	Bundles of mutually reinforcing practices explain performance — not individual practices.
Shah & Ward (2007)	Built the field's measurement instruments for lean exclusively on bundle logic.

This is the strongest available explanation for the failure rate of tool-led lean programmes. It also explains something counter-intuitive: a partial implementation can leave an organisation *worse* off than no implementation, because a practice adopted without its complements imposes its costs without delivering its returns.

SIX

Lean in the language of capital

Operational improvement that cannot be stated in financial terms does not survive the budget cycle. Fortunately, the translation is an identity rather than an argument.

Little's Law

$$\text{WIP} = \text{Throughput} \times \text{Lead Time}$$

Little (1961)

Hold throughput constant and the identity says something immediate: **cutting lead time cuts work in process proportionally.** And work in process is not an operational abstraction — it is inventory, and inventory is invested capital.

`lead time ↓` → `WIP ↓` → `inventory ↓` → `invested capital ↓` → `ROIC ↑`

Flow improvement is a balance-sheet intervention that merely happens to take place on a shop floor.

IN PROJECT TERMS

Little's Law holds for a portfolio exactly as it holds for a production line, and the portfolio version is the more useful one. Take projects in flight as WIP, projects completed per year as throughput, and the identity returns the average time a project spends in the organisation:

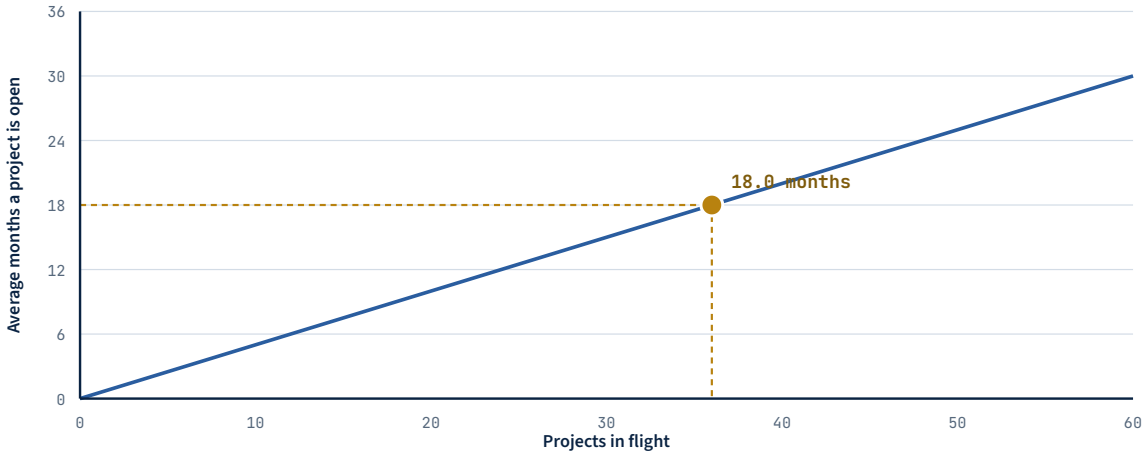
`average project lead time = projects in flight ÷ projects completed per year`

The arithmetic is unforgiving. An organisation running 20 projects and finishing 10 a year has a two-year average lead time, whatever its schedules say. Halving the number in flight halves the lead time with the same people and the same budget — and throughput usually *rises*, because the context switching that consumed the difference disappears.

This is the counter-intuitive law every steering committee eventually rediscovers at cost: **starting less finishes more.**

Little's Law, as a portfolio instrument

Shown here at its default setting. The version on the course page has sliders, and the chart and the sentence below move with them.



With **36** projects in flight and **24** finished a year, the average project is open for **18.0 months**. Cutting the number in flight to 22 — with no change to capacity, no new tool and nobody working faster — brings that to **11.0 months**.

Figure 5. Little's Law applied to a portfolio: the average time a project stays open is the number in flight divided by the number finished per year. Throughput is set by capacity and is hard to move quickly; the number in flight is a decision somebody takes every week. The line is straight because the law is an identity, not a forecast — it holds on averages, for a system that is stable and where what starts eventually finishes.

READING THE CHART HONESTLY

The relationship is an identity, so there is nothing to forecast and nothing to dispute. What it does require is a stable system in which what starts eventually finishes: a portfolio that quietly abandons a third of what it begins is not described by this line, and the first useful thing such an organisation can do is count the abandonments.

WE DO THIS CALCULATION LIVE IN THE SESSION

A. Projects currently in flight in your organisation or unit. Count everything authorised and not yet closed, including the ones nobody has touched this month. Those count — they are WIP.
B. Projects completed and closed in the last twelve months.

A ÷ B is your real average project lead time. Approximate figures are enough. The interesting number is not the result — it is the gap between that result and the lead time your organisation believes it has.

Where lean shows up in return on invested capital

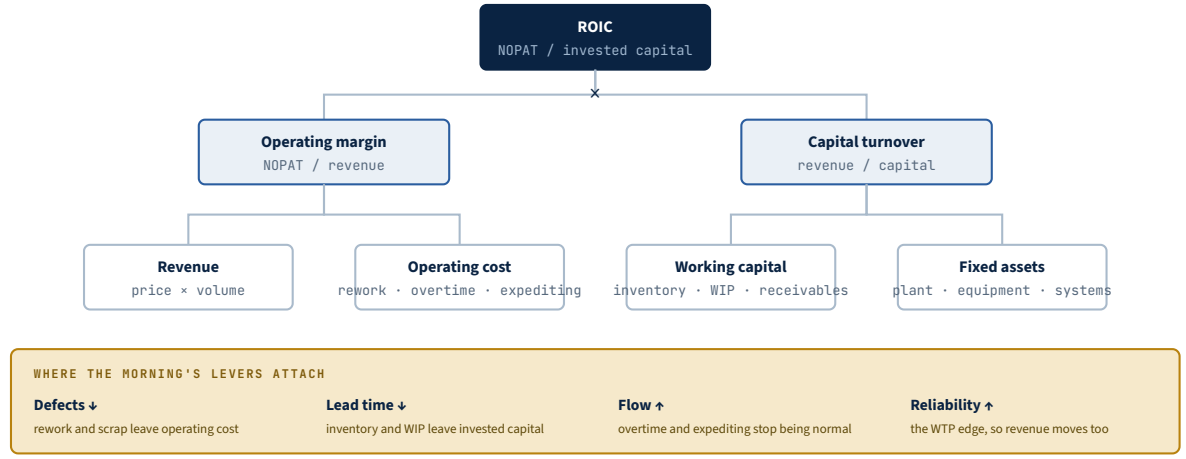


Figure 6. Return on invested capital, decomposed. The identity has two branches and the morning's work lands on both: removing defects and expediting lifts the left branch, and removing lead time lifts the right one by taking capital out of the denominator. A programme argued only as cost saving is claiming half of what it does.

ROIC is NOPAT divided by invested capital, and it splits into two independent levers:

Margin — the numerator

NOPAT ÷ revenue. Lean moves it through scrap, rework, yield and equipment effectiveness.

This is where every standard lean business case is written.

Capital turnover — the denominator

Revenue ÷ invested capital. Lean moves it through inventory, footprint, capacity buffers and the cash cycle.

This is where nobody in the operations meeting is looking.

A worked scenario

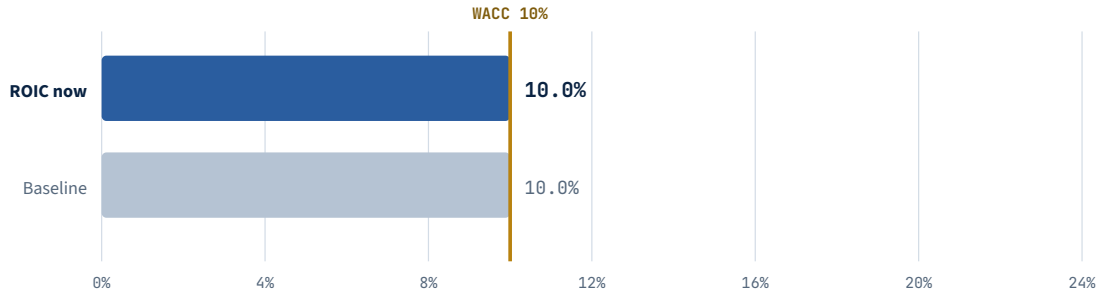
Lead time is cut from 20 days to 8. Defect reduction adds 20 to NOPAT. Revenue is unchanged.

Measure	Before	After
Inventory	300	120
Invested capital	800	620
Operating margin	8%	10%
ROIC	10%	16.1%
WACC	10%	10%
Economic profit	0	38

Margin improved by 25%. Capital turnover improved by 29%. **The denominator did more work than the numerator** — and the denominator is the half of the equation that operational business cases almost never mention. A firm that was exactly earning its cost of capital is now creating economic profit, without a single additional unit of revenue.

What an operational improvement does to return on capital

Shown here at its default setting. The version on the course page has sliders, and the chart and the sentence below move with them.



INVENTORY	INVESTED CAPITAL	NOPAT	ROIC	ECONOMIC PROFIT
300	800	80	10.0%	+0

> Return equals the cost of capital exactly. The accounting profit is real and the economic profit is zero.

Figure 7. The same improvement, read as a return on capital. Shortening lead time releases inventory, which leaves the denominator; removing rework and expediting lifts the margin, which raises the numerator. The grey bar is the starting position, where return exactly equals the cost of capital and the business creates no economic value at all. The gold line is the only threshold that matters.

The two inputs are the morning's two mechanisms, and the figure is worth operating at least once with a deliberately unhelpful setting: raise the margin as far as it goes while leaving lead time at thirty days, and watch how much of the improvement the denominator eats. That is the shape of a cost-reduction programme that never touched flow.

Both pillars move both edges of value

Mechanism	Effect
Jidoka → cost	Defects removed at source; the cost of poor quality collapses. The classic business case.
Jidoka → WTP	Zero-defect delivery lets a B2B customer eliminate incoming inspection and cut safety stock. A willingness-to-pay argument with a hard number attached.
JIT → cost	Less inventory, handling, expediting and obsolescence; a smaller footprint.

JIT → WTP	Customers pay for speed and predictability, because short lead times cut their own forecasting risk.

SEVEN

DMAIC — conducting improvement, analysing problems

Everything so far describes what good looks like. DMAIC is the method for getting there from where you are — a discipline for conducting an improvement and, above all, for analysing a problem before paying to solve it.

The failure mode it exists to prevent

Six Sigma was developed at Motorola in the 1980s and scaled across industry in the following decade. Its statistical ambition — a process so consistent that specification limits sit six standard deviations from the mean, roughly 3.4 defects per million opportunities — is the part everyone remembers. It is not the part that matters most to a senior manager.

What matters is the discipline. An experienced manager recognises a problem in seconds and already knows what to do. That recognition is pattern-matching against previous problems: fast, frequently right, and entirely unverified. When it is wrong, the organisation finds out after it has paid for the solution — and because the supposed cause was never written down, the error teaches nothing.

DMAIC does not exist to distrust experience. It exists so that the experienced manager's hypothesis is written down before it is funded, and can therefore be refuted while refutation is still cheap.

The five phases

Phase	The question it answers	Representative tools
Define	What is the problem, whose is it, what does it cost, and where does the scope end?	Project charter · SIPOC · voice of the customer
Measure	What is today's performance, measured in a way anyone would trust?	Operational definitions · measurement system analysis
Analyze	What is the cause — verified, not assumed?	Five whys · Ishikawa · Pareto · hypothesis testing
Improve	Does the solution work in a pilot, before it is scaled?	Controlled pilot · poka-yoke · design of experiments
Control	Does the gain survive the dissolution of the project team?	Control charts · control plan · standardised work

Two observations that are easy to skip. First, Define and Measure look like bureaucracy and are half the work: a badly defined problem has no verifiable solution, and a measure nobody trusts ends every discussion in opinion. Second, **poka-yoke appears in this table and also in the Jidoka pillar above.** Lean and Six Sigma are not rival systems; they share mechanisms and attack different enemies — flow and waste on one side, variation and defects on the other.

Analyze, in practice

This is the phase we will work through live in the session, using one of your own pre-work sentences.

The five whys is the simplest root-cause technique and the easiest to perform badly. Two rules govern it:

- **A cause attributed to a person is not a process cause.** "Because he did not deliver on time" is a restatement of the symptom. The next question is always: what in the process allows that to happen without anyone noticing?
- **Stop when you reach something somebody controls.** A chain that bottoms out in "because the market is volatile" has gone one step too far; a chain that stops at the second link has not gone far enough.

What the technique reliably demonstrates is that the cause people name first is usually the third or fourth link in the chain — that is, still a symptom. Pareto analysis then answers a different question: of the many verified causes, which few account for most of the effect.

Where initiatives die

In Analyze

The team arrives with its preferred solution and uses the phase to justify it. The diagnostic symptom is simple: the root cause turns out to be whatever someone already wanted to buy.

In Control

The gain exists for as long as the team exists. Once the team disperses, the process regresses within months, because nobody inherited the standard.

The cheapest safeguard against both is a **tollgate** at the end of each phase, chaired by the executive sponsor who owns the business result rather than by the specialist running the project. Whoever conducts the work is not whoever approves it.

IN PROJECT TERMS

A DMAIC initiative is itself a project, and the vocabulary overlaps almost exactly: the charter is a project charter, the tollgates are stage gates, the sponsor is the sponsor. What differs is where the weight sits. A conventional project plan commits to a solution at authorisation; DMAIC forbids naming the solution until Analyze has verified a cause. The gate between Analyze and Improve is the one with no equivalent in most governance frameworks, and it is the one that does the work.

Control maps onto benefits realisation and handover to operations — the phase most project methodologies describe and fewest organisations actually perform. A project formally closed with the process owner untrained and the standard unwritten has not transferred anything. It has only stopped.

Where the method does not apply

DMAIC assumes a repeatable process with a definable defect. Applied to the exploratory front end of a business — research, invention, early-stage product discovery — it suppresses precisely the variation that discovery depends on. The organisation needs two regimes, one for exploring and one for exploiting, and clarity about which process is in which.

EIGHT

The third equation

The identity was stated in Part Two, before there was anything to test it against. Everything since has been that test: the waste, the queues, the variation, the system of complementarities, the method for improving it. This part puts the three terms back together.

$$V = C \times (A + P)$$

Lima (2026), Ch. 1 — after Van Mieghem & Allon (2015)

Term	What it is	Visibility
A — Assets	Plants, machines, systems. What the pillars stand on.	Visible and purchasable
P — Processes	How work flows across the assets. The pillars and the foundation.	Partly visible, imitable with effort
C — Capability	The repeated ability to improve P. The centre of the house.	Barely visible, rarely imitable

The form of the equation is the claim. **Assets and processes are additive and purchasable; capability is a coefficient.** You can buy a plant, license a platform, hire a consultant to redraw a process. You cannot buy the organisation's repeated ability to improve. And if C approaches zero, V approaches zero regardless of what A and P are worth.

IN PROJECT TERMS

Sort a PMO's assets by the equation and the result is uncomfortable. The scheduling tool, the portfolio platform, the dashboards: **A**. The methodology, the templates, the gate model, the certified staff: **P**. Both are purchasable, and both are available to every competitor at roughly the same price.

C is what is left: whether this organisation reliably gets better at delivering than it was two years ago, and whether that improvement survives the departure of the people who drove it. It is also the reason two organisations running the same methodology, with the same certifications and the same tooling, deliver at visibly different standards — a gap no maturity assessment fully explains and every experienced manager recognises.

Which raises the question the session closes on, and which the afternoon takes up directly: if C is the multiplier, what actually builds it?

The equations, in one place

Equation	What it establishes
Value stick WTP − SOC	Lean pushes both edges outward at once: Jidoka and JIT each lower cost and raise willingness to pay.
Competitive advantage $VC = (WTP_A - C_A) - (WTP_B - C_B)$	The comparison form of the first. Creating value is necessary; creating more of it than the rival is what the firm is paid for. When an improvement is copyable, the two brackets converge and VC returns to zero.
Economic profit ROIC vs. WACC	Lean lifts the numerator and cuts the denominator, moving return above the cost of capital.
VCAP $V = C \times (A + P)$	It does so only for organisations that build the multiplier rather than buying the bracket.

TWO IDEAS TO BRING TO THE SESSION

Porter's objection. Operational effectiveness is imitable, so its gains compete away; strategy-grade advantage must live somewhere harder to copy. Where, exactly — and what does that imply for an improvement programme whose entire output is better processes?

Spear and Bowen's finding. What distinguishes Toyota is not its tools but its *rules-in-use*: tacit rules governing how flow paths are specified and how every improvement is run as a controlled experiment. Tacit rules can be observed and still not transferred. Hold that thought against the puzzle in Part One.

NINE

Group activity — The Value Case

The final block of the session is worked in groups. You will take one real project and appraise it through the three lenses built in these notes: where it moves willingness to pay, where it removes waste, and what it actually builds.

> **THE RAW MATERIAL**
>
> **One authorised project from your own organisation** — in flight or recently approved — that you know well enough to defend. It does not need to be large, successful, or yours. Groups choose one project from among those their members propose, and the choice takes ninety seconds: pick the one the group disagrees about, not the one with the best business case.

Why this appraisal differs from a standard business case

A conventional business case answers one question: does the return exceed the cost? That question is necessary and it is not sufficient, because it is silent on three things that decide whether the return survives.

- **Which edge of value does this move?** A project that only lowers cost is competing on the edge where every rival is also working. A project that raises willingness to pay is harder, rarer, and usually under-argued because the number is harder to produce.
- **Is the cost saving a named waste, or an aspiration?** "Improves efficiency by 15%" is a target. "Eliminates the three re-approvals caused by incomplete submissions" is a waste with a name, and only the second can be verified afterwards.
- **Does the organisation end up more capable, or only better equipped?** A project that delivers assets and processes can be replicated by any competitor willing to spend the same money. A project that leaves the organisation better at improving cannot.

Those three questions are one question in disguise, and VC states it: a project is worth defending to the extent that it widens the gap between *(WTP − cost)* for your firm and the same quantity for the competitor you are actually losing deals to. Field 6 of the canvas is where that is settled.

The Value Case Canvas

Groups of three or four · 15 minutes to complete · 3 minutes to defend

The group picks **one** of the projects brought by its members and completes six fields, one of which is a short value stream sketch. Disagreement inside the group is the point — if you all agree in the first three minutes, you have chosen a project that is too easy.

Fill this in here. One person per group types; everything is saved in that browser as you go, and the buttons at the end produce a file to send on.

1 · The project

What it is, who sponsors it, and what it was approved to achieve.

One line

2 · The WTP edge — does the customer notice?

Identify whether this project raises what someone would be willing to pay, and be precise about the mechanism.

Competitive priority affected

- [ ] Cost
- [ ] Quality
- [ ] Delivery speed
- [ ] Flexibility
- [ ] None — this project does not move any of them

Quality dimension, if any

Conformance lowers cost. Reliability, durability and serviceability raise willingness to pay.

- [ ] Performance
- [ ] Features
- [ ] Reliability
- [ ] Conformance
- [ ] Durability
- [ ] Serviceability
- [ ] Aesthetics
- [ ] Perceived quality

Order winner or qualifier?

- [ ] Order winner — it produces preference
- [ ] Qualifier — its absence eliminates us, its presence attracts nobody
- [ ] Neither

Who is the customer here — and would they notice?

If the honest answer is that nobody outside would notice, say so. Many sound projects move only the cost edge.

The customer, and yes or no

3 · Map the stream — where the time actually goes

Sketch the process the project touches, at the level of three to six steps. You are not drawing a full value stream map in ten minutes; you are producing the two numbers a map exists to produce, and locating the gap between them.

The steps, in sequence

Three to six, each with who performs it. Separate them with arrows: request → triage → approval → execution.

step (who) → step (who) → step (who)

Process time — the time anyone is actually working on it

Add up the hands-on time across the steps. An estimate the group can defend is enough.

e.g. about 4 hours

Lead time — elapsed time from first step to last

Calendar time, not working time. The ratio of the two is your flow efficiency, and it is usually between two and ten per cent.

e.g. about 3 weeks

Where the gap sits

Name the one or two places that hold most of the elapsed time. That is where the project should be aimed, and it is frequently not where it is aimed today.

The queue, the approval, the handoff

4 · The cost edge — name the waste

Now name which of the wastes lives in the gap you just located. Name it from the list; do not substitute a benefit for a waste.

Primary waste removed

- [ ] Overproduction — the root waste
- [ ] Waiting
- [ ] Transport
- [ ] Over-processing
- [ ] Inventory
- [ ] Motion
- [ ] Defects
- [ ] Unused talent
- [ ] None of the eight — see note below

Where in the process it occurs today

The step or the queue

Secondary wastes affected

Optional.

Optional

Order of magnitude, if you can estimate it

Optional. Days, hours, percentage — whatever you can defend.

Optional

The test: "reduces costs" and "improves productivity" are not wastes — they are consequences. If the group cannot name a waste from the list, the project may be removing something real that nobody has yet defined, which is itself a finding worth reporting. And if the waste removed is *overproduction*, note it: that is the root waste, and the project is probably worth more than its business case claims.

5 · What the project actually builds

Decompose the deliverable into the three terms. Most projects touch more than one; few touch all three.

A — Assets

What is bought or built?

P — Processes

What flow changes, and how?

C — Capability

What will this organisation be able to do afterwards that it could not do before — and that does not depend on the new system?

The test for C is strict. Training people to operate the new system is not capability — it is part of delivering the asset. Capability means the organisation ends the project *better at improving processes than it was before*: a routine that survives, a standard that is now owned, a group that can now run the next improvement without external help. Most projects leave C untouched, and that is a legitimate answer.

6 · The verdict

How long before a competitor could replicate this?

- ☐ Months
- ☐ A year or two
- ☐ Years
- ☐ They could not — and here is why

The decision

- ☐ Fund
- ☐ Fund with conditions
- ☐ Do not fund

If conditions: what would have to be added to build C?

One line

The group presents field 1 and field 6, with the reasoning that connects them. Three minutes, one spokesperson.

What to expect

Two results are common enough to be worth naming in advance, and neither is a failure of the project you chose.

Only one edge moves

Most project portfolios are built almost entirely of cost-edge projects. That is not wrong, but it is a choice, and usually an unexamined one. The question it raises: what would it take for this same project to also move willingness to pay — and would that cost much more?

C is empty

Projects deliver assets and processes because those are what can be specified, scheduled and signed off. Capability resists all three. The question it raises: if A and P are purchasable by any competitor, what in this portfolio is actually building advantage?

Bring the completed canvas with you to the afternoon session. It is the input to the module's Execution Optimization Challenge.

TEN

Going further — how capability accumulates

This section is optional. The session closes on the claim that C is the multiplier and that it cannot be bought. The obvious next question — then how is it built? — is a research topic in its own right, and what follows is the short version of what the evidence supports.

Repetition alone does not accumulate capability

The learning curve — unit cost falling with cumulative output — has been documented since Wright's 1936 study of aircraft manufacturing. But the later literature found it is not a uniform law. Argote and Epple (1990) showed that organisations vary widely in how fast they learn, and attributed the variation to organisational forgetting, personnel turnover, knowledge transfer from elsewhere, and scale.

That finding moves the question. If accumulated experience does not determine learning speed, and organisations with comparable experience learn at different rates, the explanation lies in **what organisations do with the experience** — the deliberate component, not the automatic one.

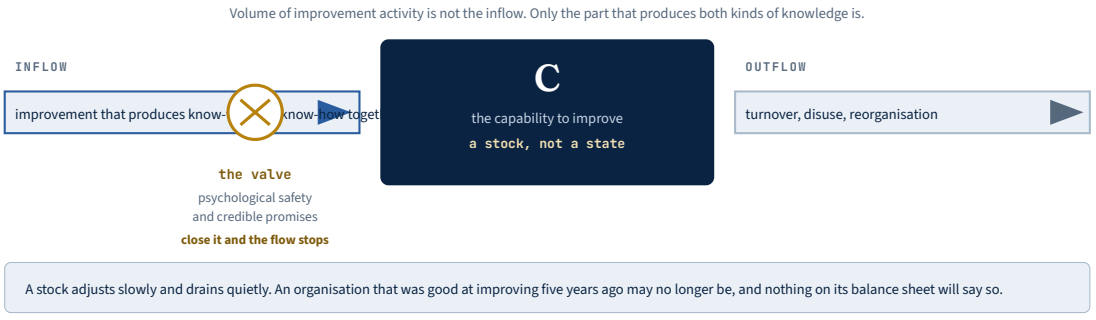


Figure 8. Capability as a stock with a regulated inflow and an unregulated outflow. The valve is the part most improvement programmes leave out: where surfacing a problem is costly, the measured performance improves and the real performance does not.

The finding worth remembering

Two studies of the same factory, a decade apart, give the sharpest decomposition available. Mukherjee and colleagues (1998), across 62 quality improvement projects, separated **conceptual learning** — understanding the cause-and-effect relationships, the *know-why* — from **operational learning** — validating the actions that produce the result, the *know-how*.

Lapré and colleagues (2000) then linked the learning activity of those projects to the factory's actual waste rate. The result is uncomfortable and precise:

> **LAPRÉ ET AL. (2000)**
>
> **Only about a quarter of the improvement projects accelerated the reduction of waste** — and they were the ones that produced know-why and know-how *simultaneously*. The remainder either slowed the reduction down or had no measurable effect at all.



Figure 9. The classification behind the finding, after Lapré, Mukherjee and Van Wassenhove (2000). Projects were sorted by whether they produced an understanding of the mechanism and whether that understanding was validated in the work. Only the cell that produced both accelerated the reduction of waste.

A project that produces know-how without know-why leaves the organisation able to repeat an action it does not understand. A project that produces know-why without know-how leaves it with a diagnosis nobody implemented. Neither accumulates anything. This is the empirical core of the DMAIC argument in Part Seven, and it explains why the Analyze phase is not a formality.

Pisano (1994) adds a boundary condition from 23 pharmaceutical process-development projects: where the underlying science is deep and well understood, learning *before* execution — laboratory experimentation — speeds development. Where it is not, as in biotechnology, it does not. The effective form of learning depends on how much is already known about the process.

WHERE THIS GOES

If improvement projects are how capability accumulates, two questions follow immediately: under what conditions do they work, and how is the accumulated stock lost again? Those are the closing questions of the *afternoon* session, after benchmarking has shown why the stock cannot be acquired from anyone else.

ELEVEN

Sources

Works cited in these notes and in the session.

- Brandenburger, A. M., & Stuart, H. W. (1996). Value-based business strategy. *Journal of Economics & Management Strategy*, 5(1), 5–24.
- George, M. L. (2002). *Lean Six Sigma: Combining Six Sigma Quality with Lean Speed*. McGraw-Hill.
- Rother, M., & Shook, J. (1999). *Learning to See: Value Stream Mapping to Add Value and Eliminate Muda*. Lean Enterprise Institute.
- Serra, C. E. M., & Kunc, M. (2015). Benefits realisation management and its influence on project success and on the execution of business strategies. *International Journal of Project Management*, 33(1), 53–66.
- Shewhart, W. A. (1931). *Economic Control of Quality of Manufactured Product*. Van Nostrand.
- Whyte, J., Stasis, A., & Lindkvist, C. (2016). Managing information and control in digitalised project environments. *Production Planning & Control*, 27(16), 1345–1354.
- Womack, J. P., & Jones, D. T. (1996). *Lean Thinking: Banish Waste and Create Wealth in Your Corporation*. Simon & Schuster.
- Garvin, D. A. (1987). Competing on the eight dimensions of quality. *Harvard Business Review*, 65(6), 101–109.
- Hill, T. (2000). *Manufacturing Strategy: Text and Cases* (3rd ed.). McGraw-Hill.
- Hopp, W. J., & Spearman, M. L. (2004). To pull or not to pull: what is the question? *Manufacturing & Service Operations Management*, 6(2), 133–148.
- Argote, L., & Epple, D. (1990). Learning curves in manufacturing. *Science*, 247(4945), 920–924.
- Hopp, W. J., & Spearman, M. L. (2011). *Factory Physics* (3rd ed.). Waveland Press.
- Lapré, M. A., Mukherjee, A. S., & Van Wassenhove, L. N. (2000). Behind the learning curve: Linking learning activities to waste reduction. *Management Science*, 46(5), 597–611.
- Lima, T. F. (2026). *Excelência Operacional: da Estratégia à Execução*. Atena Editora. Open access.
- Lima, T. F. *O que não se copia: capacidade, práticas e desempenho em operações*. Manuscript in preparation.
- Mukherjee, A. S., Lapré, M. A., & Van Wassenhove, L. N. (1998). Knowledge driven quality improvement. *Management Science*, 44(11-part-2), S35–S49.
- Pisano, G. P. (1994). Knowledge, integration, and the locus of learning. *Strategic Management Journal*, 15(S1), 85–100.
- Linderman, K., Schroeder, R. G., Zaheer, S., & Choo, A. S. (2003). Six Sigma: a goal-theoretic perspective. *Journal of Operations Management*, 21(2), 193–203.
- Little, J. D. C. (1961). A proof for the queuing formula L = λW. *Operations Research*, 9(3), 383–387.
- MacDuffie, J. P. (1995). Human resource bundles and manufacturing performance. *ILR Review*, 48(2), 197–221.
- Milgrom, P., & Roberts, J. (1990). The economics of modern manufacturing. *American Economic Review*, 80(3), 511–528.
- Modig, N., & Åhlström, P. (2012). *This Is Lean: Resolving the Efficiency Paradox*. Rheologica.
- Ohno, T. (1988). *Toyota Production System: Beyond Large-Scale Production*. Productivity Press.

- Repenning, N. P., & Sterman, J. D. (2001). Nobody ever gets credit for fixing problems that never happened. *California Management Review*, 43(4), 64–88.
- Rother, M., & Shook, J. (2003). *Learning to See*. Lean Enterprise Institute.
- Schroeder, R. G., Linderman, K., Liedtke, C., & Choo, A. S. (2008). Six Sigma: definition and underlying theory. *Journal of Operations Management*, 26(4), 536–554.
- Shah, R., & Ward, P. T. (2007). Defining and developing measures of lean production. *Journal of Operations Management*, 25(4), 785–805.
- Spear, S., & Bowen, H. K. (1999). Decoding the DNA of the Toyota Production System. *Harvard Business Review*, 77(5), 96–106.
- Van Mieghem, J. A., & Allon, G. (2015). *Operations Strategy: Principles and Practice* (2nd ed.). Dynamic Ideas.

International Certificate in Project Management and Organizational Performance

Module 3 — Methods, Tools, and Operational Efficiency · Session 1 reading notes · 17 October 2026

Sorbonne Business School · IDGP Academy · Prof. Dr. Thiago Lima